

# 2024 High Performance Computing Coursework

Amin Donald  
University of Bristol  
Bristol, England  
sq21386@bristol.ac.uk

**Abstract**— In this paper, we attempt to complete the University of Bristol 2024 High Performance Computing coursework. This entails optimising some existing code for a Lattice Boltzmann CFD simulation on the BlueCrystal Phase 4 supercomputer.

**Index terms**—*High Performance Computing*

## I. SERIAL OPTIMISATION

### A. Compiler Arguments

For gcc, the important arguments we used were -Ofast and -march=native. -Ofast tells the compiler to use all the optimising techniques in the -O3 argument, but also enables some techniques that, while beneficial, can cause UB in some programs (e.g -fallow-store-data-races). The -march=native argument (and similarly the -xHost arg for icc) tells the compiler to not only tune its optimisations to the specific kind of hardware we’re compiling for, but to try and use platform specific instructions, such as vector instructions.

### B. Loop Fusion

To begin with, the code calls 5 functions in a given timestep: accelerate, propagate, rebound, collision and avg\_velocity. There is also a cells array, the main simulation space, and a tmp\_cells array, a scratch space for calculation. Each of these functions run their own loop, so the first optimisation we did was to make timestep only have one loop. After manually inlining each function and removing each’s inner loop, it becomes apparent that a cell only collides or rebounds, but never both, so we made them conditional, and that the collision step already calculates the average velocity, so we just return this value from the timestep function. Another optimisation is the “pointer swap”. This is where, instead of writing back to the cells array in the collide and rebound step, we only read from cells, write back to tmp\_cells, and then just swap their pointers after each timestep. This means we only mutate one array rather than two, leading to better memory use and caching.

## II. VECTORISATION

### A. Compiler Arguments

We used icc to vectorise the code, as it was well documented. Other than -xHost so that the compiler knows the cpu can run vector instructions, the directives to vectorise exist within the source code. Though, -opt-report can tell you to what extent your code was vectorised.

### B. Structure of Arrays

Instead of the previous array of structs containing cell speeds, we switched to a struct of arrays, where the first array contained all the speeds in one direction, the second contained all the speeds in another direction, and so on. This provides our compiler with memory that is contiguous for one parameter (i.e. all the speeds in some direction are one after the other), and considering each timestep does the same operation, this lets the compiler vectorise.

### C. Memory alignment

To use vector instructions, the “vector” in memory must be aligned. In our case we used icc’s \_mm\_malloc function to ensure the cell arrays were aligned in the physical memory, and then within the timestep loop, used the \_\_assume\_aligned function to tell the compiler that the data is aligned for the purposes of optimisation. We also use an \_\_assume function to tell the compiler that the loop iterations are divisible by 64, so that it can do several iterations in one vectorised iteration, and not need to worry about having to do a partial iteration.

## III. PARALLELISATION

### A. Compiler Arguments

We used OpenMP to parallelise our code, so we need to tell icc we are using it with the -qopenmp argument.

```

1  ...
2  #pragma omp parallel for collapse(2)
    reduction(+ : tot_u) reduction(+ :
    tot_cells)
3  for (int jj = 0; jj < params.ny; jj++) {
4      for (int ii = 0; ii < params.nx; ii++) {
5          ...

```

Listing 1: Excerpt from the parallelised timestep loop

#### B. “parallel for” pragma

OpenMP contains very versatile pragmas that can, in just one or two lines, parallelise code. In our case, on the main loop we use the pragma shown in Listing 1. `omp parallel`, and `for` read literally as “parallelise this for loop”, but by way of splitting all the iterations among the different threads. `collapse(2)` tells OpenMP that the `for` loop is actually two `for` loops, with one nested in the other, and that the pragma should turn them into just one `for` loop, so that the allocation of iterations can be more fine grained than just the outer loop. The reduction arguments are to do with getting a result from something that is common among all threads. None of the threads ever modify the same cell, but every thread needs to increment the number of cells and contribute to the average velocity. So `reduction(+ : tot_cells)` says “make the `tot_cells` variable local to each thread so there are no store race-conditions, but then after the loop is done, add each thread’s `tot_cells` together so that we end up with one total `tot_cells` variable”, and similarly for `tot_u`.

## IV. PERFORMANCE ANALYSIS

Size	Code kind	runtime (s)
128x128	Serial	21.679
	Vector	7.243
	Parallel 28 thread	1.822
256x256	Serial	175.305
	Vector	66.706
	Parallel 28 thread	20.516
1024x1024	Serial	729.734
	Vector	336.093
	Parallel 28 thread	84.773

Figure 1: Runtimes of different code instances on different data. The parallel data is averaged over 4 runs.

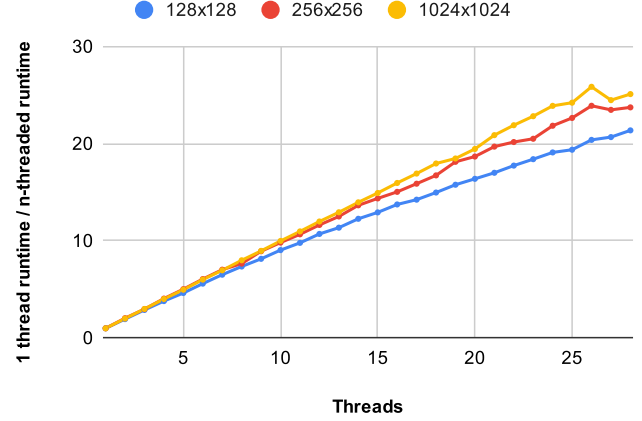


Figure 2: Parallel performance plotted against number of threads. Data is averaged over 4 runs.

As somewhat expected, the parallel perf is faster than the serial perf (Figure 1). The trend that the parallel perf has (Figure 2) is quite linear, which seems good, but as the code does not meet the assignment’s benchmark, this is linear performance scaling with poor runtimes, and so likely doesn’t matter.